The Demobilizing Effect of Selective Exposure and Heterogeneous Responses to Violence*

Carolina Torreblanca†

March 22, 2026

Abstract

Whether those who bear the burden of violence stay in political life or retreat shapes the government's cost of failing to provide security. Yet much of what we know about the participatory consequences of violence comes from populations unlikely to experience it. I show that when exposure is selective and reactions heterogeneous, population-average effects and effects among the exposed can reverse in sign. I develop a principal stratification framework encoding both features, and a partial identification procedure applicable to any binary treatment and binary outcome. Applied to criminal victimization across 94 surveys in 20 Latin American countries, population-average effects favor mobilization but, contrary to extant findings, effects among the victimized are predominantly negative. Panel evidence from Mexico indicates those likely to demobilize are also more exposed. For treatments that are not universal and that we would never seek to universalize, the causal contrast of first-order importance is among the treated.

Word count: 9595 words

*I thank Guilherme Duarte, Guy Grossman, Kun Heo, Nicolás Idrobo, Dorothy Kronick, Giacomo Lemoli, Beatriz Magaloni, Justin Melnick, Alexis Palmer, Pablo Querubín, Arturas Rozenas, Cyrus Samii, Tara Slough, David Stasavage, Elisa Wirsching, Yuhan Zheng, audiences at NEWEPS 24, APSA 2024, and at the University of Pennsylvania's Comparative Politics Workshop for excellent comments and suggestions.

†Postdoctoral Fellow, PDRI-DevLab, University of Pennsylvania, catba@sas.upenn.edu.

Empirical social science has converged on the view that treatment effects that average across all units should serve as the benchmark for causal knowledge and policy evaluation. For treatments that could, in principle, reach anyone—vaccination campaigns, cash transfers, educational curriculum changes—population-wide causal contrasts are maximally informative: they reflect what we should expect, on average, at scale.

However, when it comes to violence, one of the treatments most studied in political science, only specific subpopulations experience it. Whether the form is criminal victimization, civil war, police contact, incarceration, or forced displacement, those exposed systematically differ from the rest of the population. In such cases, the population average answers a question about a universal exposure counterfactual that would never occur and that no one would wish for. I argue that for treatments where exposure is selective and responses are heterogeneous, and where the treatment is not one we would ever seek to universalize, the empirically relevant causal contrast is instead the effect among those exposed, not the population-wide average.

The difference between what these two causal contrasts can tell us becomes especially clear when examining the participatory consequences of violence, where exposure is not only selective, but responses are also heterogeneous, and the characteristics that determine who is exposed also shape how they respond. Participatory responses to violence range from mobilization through grievance and collective action (Blattman, 2009; Bateson, 2012) to demobilization through fear and withdrawal (Ley, 2018; Skogan, 1986). Exposure to different types of violence is routinely selective: criminals target vulnerabilities (Browning, Pinchak and Calder, 2021), police discriminate (Knox, Lowe and Mummolo, 2020), and civil war violence follows territorial logic (Kalyvas et al., 2006). The characteristics that predict political participation, such as risk tolerance, mobility, and community embeddedness (Kam, 2012; Gerber et al., 2011), overlap substantially with those that predict victimization.

To characterize when and how the population-average effect diverges from the effect among the exposed when selective exposure and heterogeneous responses coincide, I develop a principal

stratification framework with binary treatment and binary outcome (Frangakis and Rubin, 2002). Departing from the standard compliance application, I define strata by individuals' potential outcomes on the outcome of interest: whether they would participate regardless of treatment, only if treated, only if untreated, or never. This captures response heterogeneity. I then extend the framework to allow the probability of exposure to vary across these response types, capturing overlapping selective exposure. I show that the population-wide average treatment effect (ATE) depends only on the relative shares of those who would respond positively and negatively to treatment, while the average effect among the treated (ATT) also depends on how exposure is distributed across these groups, weighting each response type by its probability of treatment.

Because the ATT weights response types by their exposure probabilities whereas the ATE does not, the two estimands can differ in magnitude and even point in opposite directions. I demonstrate that whether sign reversal occurs depends on two quantities: the intensity of selection, or whether each treatment-sensitive response type is more likely to be exposed than the other, and the degree of response heterogeneity, or the relative prevalence of those who would respond positively versus negatively in the population. When exposure is sufficiently skewed toward those who would respond in the opposite direction of the population average, a positive population-average effect masks a negative effect among the exposed.

The framework establishes that sign reversal occurs when the response type that is less prevalent in the population is sufficiently overrepresented among the exposed. I show that for a given estimated population-average effect and a prior on the intensity of differential selection, the minimum population share of this countervailing type required to flip the sign of the effect among the exposed is a closed-form expression that requires no empirical input. I provide a lookup table (Table 2) that reports these thresholds across a range of ATE magnitudes and selection ratios, applicable to any setting with binary treatment and binary outcome. Applied researchers can use this table to assess what they must believe about selective exposure and stratum sizes for their population-average and treated-average effects to agree in sign.

I apply the framework to reassess the participatory consequences of criminal victimization in Latin America, the world's most violent region (Vilalta, 2020). Extant studies have identified a positive population-average effect of victimization on non-electoral political participation (Bateson, 2012; Dorff, 2017; Boulding, Mullenax and Schauer, 2022), a finding that has been read as evidence that victimization may foster a more engaged citizenry. I investigate whether the conclusion holds when the causal contrast is defined over those actually victimized.

Using the identities derived from the framework, I develop a partial identification procedure that characterizes the full set of ATEs and ATTs consistent with four moments observable in data: the prevalence of victimization, the prevalence of participation, and the participation rates among the victimized and non-victimized. I extract these moments from 94 waves of the LAPOP survey, covering approximately 150,000 respondents across 20 Latin American countries (LAPOP, 2022), an extended version of the Latin American data used in Bateson (2012). Each combination of response-type shares and exposure probabilities consistent with these moments defines a parameter vector, yielding a specific ATE, ATT, and implied degree of selective exposure.

Across four measures of non-electoral participation, the majority of parameter vectors consistent with the data yield a negative effect among the victimized: between 53% and 93%, depending on the outcome. The population-average effect, by contrast, is positive in 63%–85% of those same vectors, consistent with the mobilization finding in the literature. Further, the two estimands disagree in sign across much of the identified set. Consequently, under a flat prior over the parameter space, victimization demobilizes those who experience it for the majority of what the data allows, even as the population average suggests mobilization.

The partial identification maps what is possible across 20 Latin American countries, and for most of the parameter space the effect among the victimized is negative. But cross-sectional data cannot tell us where in that space we actually are. For that, I turn to a quarterly rotating panel survey of approximately 174,000 Mexican city residents surveyed over 27 quarters (INEGI, 2024). Between-individual comparisons show that soon-to-be victims are more mobile, witness more

criminal and gang activity, and report more violent conflicts than their neighbors before being victimized. Conversely, within-individual comparisons show that after victimization, individuals become less mobile, less trusting of law enforcement, and more fearful than before, consistent with demobilization.

Taken together, the panel evidence tells us that demobilizers are more likely to be exposed to crime than mobilizers. But differential exposure alone does not sign the ATT. The framework tells us that the sign also depends on whether mobilizers are sufficiently more numerous in the population to overcompensate for their lower exposure. The population-average effect is the difference between the share of mobilizers and the share of demobilizers in the population. When that difference is small, the two groups are of comparable size, and even moderate differential exposure is enough to reverse the sign. Since the population-average effects in the literature are substantively small, on the order of 1 to 4 percentage points (Bateson, 2012), mobilizers are unlikely to outnumber demobilizers by a large enough margin. Given the directional evidence from the ENSU panel and the small magnitude of published ATEs, the effect among the victimized is signed negative for most non-electoral participation outcomes.

The participatory consequences of violence have been studied not only for crime but also for civil war and policing, and in each case exposure is plausibly selective and responses heterogeneous. I apply the partial identification procedure to data from Blattman (2009) on forced abduction by the LRA in northern Uganda and from Walker (2020) on direct police contact in the United States, recovering the identified sets of ATEs and ATTs for electoral and non-electoral participation in each setting. In both cases, the identified sets are consistent with the same structure: population-average effects favor mobilization while effects among the exposed favor demobilization, though without panel evidence these extensions cannot constrain the identified sets with the same force.

The paper makes three main contributions. First, the paper contributes to the literature on the political consequences of violence. Influential studies find that criminal victimization (Bateson, 2012; Dorff, 2017; Boulding, Mullenax and Schauer, 2022), civil war exposure (Blattman, 2009;

Bellows and Miguel, 2009), and police contact (Walker, 2020) increase political participation, though the evidence across settings is mixed. I show that across all three settings, the mobilization finding reflects population-average effects that mask demobilization among those exposed, since the people most likely to experience violence are often also the people most likely to withdraw from political life after experiencing it.

Second, it contributes to work on what researchers learn from different estimands under selective exposure. Heckman and Vytlacil (2005) showed that no single treatment effect parameter summarizes an intervention when selection into treatment covaries with how individuals respond to it, using a fully specified structural model. Samii, Paler and Daly (2017) proposed targeting estimands defined over the populations most relevant for policy. I build on both insights but depart in two ways. First, I derive a closed-form condition under which the ATE and ATT disagree in sign—a condition that depends on only two interpretable quantities, the intensity of differential selection and the relative prevalence of response types. This yields a practical diagnostic (Table 2) that any applied researcher can use without running the partial identification procedure. Second, the partial identification recovers the joint identified set of both estimands from four observable moments, making visible not just whether sign reversal is possible but how much selective exposure is required to produce it. The framework requires no structural model and no assumptions beyond the existence of principal strata.

Third, it contributes to research on the accumulation of causal evidence by showing that population-specific estimands incorporate heterogeneity differently, and that this matters for whether findings across studies are commensurable. Researchers have emphasized that empirical estimates must be linked to theoretical claims through an explicit estimand (Lundberg, Johnson and Stewart, 2021) and that for findings to accumulate across studies, estimates must target the same causal quantity (Slough and Tyson, 2023). I show that under selective exposure the estimand simultaneously defines the link between theory and evidence and the population over which effects are averaged (Mesquita and Tyson, 2020). Research designs targeting different estimands are not commensu-

rable even when they examine the same treatment, outcome, and data. The finding that violence mobilizes and the finding that violence demobilizes can both be correct; they answer different questions, defined over different populations.

1 A Principal Strata Framework

In this section, I develop a principal stratification framework that accounts for both heterogeneous responses and selective treatment exposure. The framework makes precise the conditions under which the population-average effect and the effect among the exposed diverge, and yields a partial identification procedure for recovering both from observable data.

Consider a population of units indexed by i and a binary treatment $Z \in 0, 1$, for example, experiencing violence or not. Denote the potential outcome for unit i under treatment assignment z as $Y_i(z)$. Following Frangakis and Rubin (2002), I classify units into principal strata (θ) defined by their joint potential outcomes $(Y_i(1), Y_i(0))$, where the outcome itself serves as the post-treatment variable with respect to which strata are defined. Thus, by definition, the individual treatment effect for each i, $ITE_i = Y_i(1) - Y_i(0)$, will be identical for all units in the same stratum.

Denote the proportion of units in each stratum π_θ, with $\sum_\theta \pi_\theta = 1$. I depart from the canonical principal-strata setup by allowing units to have stratum-specific probabilities of receiving treatment, ρ_θ. Table 1 synthesizes the setup for the case of a binary outcome $Y_i(z) \in 0, 1$, such as whether an individual participates in a given political activity or not.

Stratum (θ)	**Share**	**Pr(Z=1)**	**Y(1)**	**Y(0)**	**ITE**
Always Participate	π_A	ρ_A	1	1	0
Participate if **T**reated	π_T	ρ_T	1	0	1
Participate if **U**ntreated	π_U	ρ_U	0	1	-1
Never Participate	π_N	ρ_N	0	0	0

Table 1: Stratum-specific individual treatment effects and treatment probabilities. Each row defines a principal stratum by its potential outcomes under treatment and control. The treatment probability ρ_θ is allowed to vary across strata.

For a binary treatment such as exposure to violence, we can classify individuals into four strata: those who always participate (A), those who never participate (N), those who participate if and only if treated (T), and those who participate if and only if untreated (U).[1]

Proposition 1. *The average treatment effect is* $ATE = \pi_T - \pi_U$.

Proposition 1 applies the principal stratification framework of Frangakis and Rubin (2002) to binary outcomes. It states that the ATE depends exclusively on the relative shares of strata T and U. This result is straightforward when we consider that all units in strata A and N contribute an ITE of zero. Consequently, the magnitude and direction of the ATE will only depend on the frequency with which T appears in the population, relative to U.[2]

Proposition 2. *The average treatment effect on the treated is*

$$ATT = \frac{\pi_T \rho_T - \pi_U \rho_U}{Pr(Z = 1)},$$

where $Pr(Z = 1) = \sum_\theta \pi_\theta \rho_\theta$.

[1] I make the stable unit treatment value assumption (SUTVA) throughout.

[2] Formal proofs for all expressions in this section are provided in Appendix Appendix A.

Conversely, the ATT, or the average expected change in outcomes for those who are treated, will be a function of population shares and also of the stratum-specific probabilities of receiving treatment. These stratum-specific probabilities scale the relative contributions of the T and U strata within the treated population. Intuitively, strata with higher treatment probabilities constitute a larger share of those observed in the treated group, and therefore should contribute more to the ATT. Further rearranging the numerator of the expression for the ATT yields an identity for the population-level impact of treatment:

$$ATT \times Pr(Z = 1) = \pi_T \rho_T - \pi_U \rho_U. \tag{1}$$

I denote this quantity ΔY, the average change in Y attributable to the realized treatment assignment.

Corollary 1. *The ATE and ATT are equal when* $\rho_T = \rho_U = Pr(Z = 1)$.

Intuitively, when $\rho_T = \rho_U = Pr(Z = 1)$, all strata face the same probability of treatment, so conditioning on being treated does not alter the relative representation of any stratum. That is, the treated group is a random draw from the population, so the ATE and the ATT average over, in expectation, identical weights.[3]

Corollary 2. *When* $\rho_T = \rho_U$*, the ATE and ATT have the same sign. Specifically,* $ATT = \frac{\rho_T}{Pr(Z=1)} \cdot ATE$.

Corollary 2 relaxes Corollary 1 by allowing the treatment-insensitive strata (A and N) to have treatment probabilities that differ from those of the treatment-sensitive strata. Because the A and N strata have a treatment effect of zero, their over- or under-representation in the treated group

[3]Corollary 1 also parallels the condition identified by Heckman and Vytlacil (2005) under which all treatment parameters coincide. They derived this result within a structural framework that models the selection process explicitly, requires an instrument, and recovers the full marginal treatment effect curve. Here, the same condition emerges from the structure of principal stratification alone, without parametric assumptions on the selection process or an exclusion restriction.

cannot change the sign of the ATT relative to the ATE. Only the relative weighting of T and U determines the sign, so $\rho_T = \rho_U$ suffices for sign agreement. The treatment-insensitive strata do, however, affect the *magnitude*: the scaling factor $\rho_T/Pr(Z = 1)$ can amplify or attenuate the ATT but cannot reverse its direction. Corollary 1 is the special case where $\rho_T = Pr(Z = 1)$, so the scaling factor equals one.

When do the ATE and ATT disagree in sign? The ATE and ATT need not agree in sign. Signs agree when composition and selection reinforce each other: for example, when the T-stratum is both larger than the U-stratum and at least as likely to be treated, or when the U-stratum is both larger and at least as likely to be treated. Signs can diverge when one stratum is larger but the other is more likely to be treated, so that the stratum that dominates the population average is under-represented among the treated. I now characterize exactly when this occurs.

To do so, I define two summary parameters. Let $\lambda = \rho_U/\rho_T$ capture differential selection between the responsive strata. When $\lambda > 1$, the U-stratum is more likely to be treated than the T-stratum; when $\lambda < 1$, the reverse holds. Last, when $\lambda = 1$ we have equal selection and necessarily sign agreement (Corollary 2). Let $\eta = \pi_T/\pi_U$ capture the relative size of the two responsive strata. Hence $\eta > 1$ when the T-stratum is larger ($ATE > 0$) and $\eta < 1$ when the U-stratum is larger ($ATE < 0$).

Proposition 3. *Let $\pi_T > 0$ and $\pi_U > 0$. Then $\text{sign}(ATT) = \text{sign}(\pi_T\rho_T - \pi_U\rho_U)$, and the ATE and ATT have opposite signs if and only if $\lambda > \eta$ when $ATE > 0$, or $\lambda < \eta$ when $ATE < 0$.*

The sign of the ATT is determined by whether the T-stratum's exposure-weighted size ($\pi_T\rho_T$) exceeds that of the U-stratum ($\pi_U\rho_U$). When $\lambda > 1$, the U-stratum is over-represented among the treated, pulling the ATT toward negative values. A sign reversal occurs when this differential selection more than compensates for the size advantage of the T-stratum, that is, when $\lambda > \eta$. The condition is symmetric: when $ATE < 0$ ($\eta < 1$), the ATT is positive only if $\lambda < \eta$. In that case, the T-stratum must be sufficiently more exposed to treatment than the U-stratum to overcome its

smaller population share.

Corollary 3. *When $ATE > 0$ and $\lambda > 1$:*

$$ATT < 0 \quad \Longleftrightarrow \quad \pi_U > \frac{ATE}{\lambda - 1}. \tag{2}$$

By symmetry, when $ATE < 0$ and $\lambda < 1$, the ATT is positive if and only if $\pi_T > |ATE| \cdot \lambda/(1-\lambda)$.[4]

For the sign of the ATT to oppose the sign of the ATE, the countervailing stratum must constitute a sufficiently large share of the total population. How large depends on two quantities: (i) the magnitude of the ATE, and (ii) how much more exposed the countervailing stratum is to treatment. When the ATE is small or λ is far from one, the threshold is low, so sign reversal requires only a modest population share of the countervailing stratum. Table 2 illustrates this relationship for a range of ATE magnitudes and selection ratios λ. The table applies to any binary treatment and binary outcome; it follows entirely from the framework and requires no empirical input.

Consider the bottom-left region of the table, where the ATE is positive and the T-stratum is also more likely to be treated ($\lambda < 1$). These cells show check marks: composition and selection reinforce each other, so the ATE and ATT always agree in sign regardless of stratum shares. Now consider the bottom-right region, where the ATE is still positive but the U-stratum is more likely to be treated ($\lambda > 1$). Here, the threshold cells show the minimum population share of the U-stratum (π_U, as a fraction of the total population) above which the ATT turns negative. For example, at $ATE = 0.04$ and $\lambda = 2$ (meaning that for every member of the T-stratum who is treated, two members of the U-stratum are treated, per capita), the ATT is negative whenever the U-stratum constitutes more than 4% of the population. Conversely, the top-left region applies when the ATE is negative and the T-stratum is more likely to be treated ($\lambda < 1$). At $ATE = -0.10$ and $\lambda = 1/3$ (meaning for every member of the U-stratum who is treated, three members of the T-stratum are

[4]Define $\lambda' = \rho_T/\rho_U = 1/\lambda$ and $\eta' = \pi_U/\pi_T = 1/\eta$. The sign-flip condition $\lambda < \eta$ is equivalent to $\lambda' > \eta'$, and the same algebra yields $\pi_T > |ATE|/(\lambda' - 1) = |ATE| \cdot \lambda/(1 - \lambda)$.

treated, per capita), the ATT is positive whenever the T-stratum exceeds 5% of the population.

		Selection ratio $\lambda = \rho_U/\rho_T$						
		T more likely treated				U more likely treated		
		$\frac{1}{3}$	$\frac{1}{2}$	$\frac{2}{3}$	1	$\frac{3}{2}$	2	3
ATE	-0.10	$\pi_T > 5\%$	$\pi_T > 10\%$	$\pi_T > 20\%$	✓	✓	✓	✓
	-0.04	$\pi_T > 2\%$	$\pi_T > 4\%$	$\pi_T > 8\%$	✓	✓	✓	✓
	-0.02	$\pi_T > 1\%$	$\pi_T > 2\%$	$\pi_T > 4\%$	✓	✓	✓	✓
	0.02	✓	✓	✓	✓	$\pi_U > 4\%$	$\pi_U > 2\%$	$\pi_U > 1\%$
	0.04	✓	✓	✓	✓	$\pi_U > 8\%$	$\pi_U > 4\%$	$\pi_U > 2\%$
	0.10	✓	✓	✓	✓	$\pi_U > 20\%$	$\pi_U > 10\%$	$\pi_U > 5\%$

Table 2: Minimum share of the countervailing stratum required for the ATT to have opposite sign from the ATE, as a function of the ATE and the selection ratio $\lambda = \rho_U/\rho_T$. When $\lambda > 1$, the U-stratum (demobilizers) is more likely to be treated than the T-stratum (mobilizers); when $\lambda < 1$, the reverse. Cells show the threshold from Corollary 3. Off-diagonal check marks indicate regions where differential selection reinforces population share differences, so the ATE and ATT always agree in sign. The $\lambda = 1$ column indicates no differential selection, so signs agree by Corollary 2. The table is symmetric around $\lambda = 1$ and $ATE = 0$.

The framework and Table 2 are entirely data-agnostic: they apply to any setting with a binary treatment and a binary outcome where treatment plausibly generates opposing responses among distinct subpopulations and exposure to treatment covaries with response type. Any empirical researcher can use the table to assess, for a given ATE magnitude and belief about differential selection, whether the ATE and ATT might diverge in sign. In the next section, I apply this framework to criminal victimization and political participation.

2 Application to Criminal Victimization

Extant theoretical accounts of the participatory consequences of criminal victimization, drawing heavily from work on civil war and policing violence, posit victimization may affect participa-

tion decisions through two competing mechanisms. The first encompasses processes that increase participation: post-traumatic growth can transform victims into civic actors, anger at victimization can generate demand for political action, shared exposure can build solidarity that lowers collective action costs, or victims may acquire instrumental incentives to demand security or punish incumbents (Bateson, 2012; Blattman, 2009; Koos and Traunmüller, 2024; Gilligan, Pasquale and Samii, 2014; Ley, 2022). The second encompasses processes that suppress it: by engendering fear, crime can raise the perceived costs of public engagement, and victimization can erode trust in the state's capacity to protect (Ley, 2018; Hale, 1996). These two broad channels, and the many mechanisms they encompass, map onto the *Participate if Treated* and *Participate if Untreated* strata, respectively.

Criminological research emphasizes that a central feature of crime is that exposure is not randomly distributed. Victimization risk varies systematically with mobility, risk preferences, socioeconomic status, and access to protection. Wealthier and more risk-averse individuals tend to be better able or more willing to avoid victimization (Skogan, 1995; Hale, 1996), whereas more mobile and risk-prone individuals embedded in their communities face greater exposure and often weaker protection (Rader, 2004; Browning, Pinchak and Calder, 2021; Boggs, 1965; Guedes, Domingos and Cardoso, 2018). These same characteristics are also correlates of political participation in the literature, with mobility, risk tolerance, and community embeddedness associated with higher rates of civic and political engagement (Gerber et al., 2011; Kam, 2012). Crime is therefore a setting in which who is exposed and how they respond are plausibly linked.

Much of the empirical literature concludes that victimization, on average, increases non-electoral civic and political engagement (Bateson, 2012; Dorff, 2017; Boulding, Mullenax and Schauer, 2022; Ley, 2022). This evidence is largely generated by research designs that target population-average or conditional average treatment effects. Bateson (2012), for example, conditions on observed covariates and invokes conditional ignorability, estimating how participation would change if victimization were randomly assigned within covariate strata. This answers a coherent counter-

factual question, but it does not describe how participation changes among those who are actually victimized when exposure is selective.[5] If subpopulations facing higher victimization risk are also more likely to be demobilized by victimization, population-average effects need not describe effects among victims.

2.1 Data

I focus on Latin America, where victimization is widespread (Vilalta, 2020; LAPOP, 2022) and where much of the research on crime and participation has been conducted (Visconti, 2020; Ley, 2018; Blanco, 2013; Marshall, 2022).

As the primary data source, I employ an extended version of the LAPOP survey, which Bateson (2012) uses for their Latin American results. I pool all surveys conducted in 20 Latin American countries between 2010 and 2019, yielding approximately 150,000 respondents from 94 survey waves. Each survey round was designed to represent the voting-age population of the country in that year.[6] LAPOP measures self-reported victimization in the past 12 months and several forms of civic and political engagement. I focus on five outcomes: protest participation, political meeting attendance, community meeting attendance, participation in community problem-solving activities, and willingness to vote in upcoming national elections if they were held next week.

2.2 Empirical Strategy

I implement the partial identification procedure using the LAPOP data. From the pooled data, I can retrieve the proportion of respondents who report being victims, $Pr[Z = 1]$ and its converse probability, the proportion of respondents who report engaging in each of the participatory activities, $Pr[Y = 1]$, and its converse, and the proportion of victimized/unvictimized respondents who report participating $Pr[Y = 1|Z = 1]$, $Pr[Y = 1|Z = 0]$, respectively. Table 3 reports the

[5]This distinction is not a critique of observational identification. Confounding concerns whether a given estimand is estimated without bias; here, different research designs can target different causal quantities for the same treatment and outcome.

[6]See Appendix A2.1.1 for sampling details.

observed probabilities in the data for each of the outcomes.[7]

Leveraging the structure of the principal-strata framework, these observed quantities restrict the possible values of the population shares of each stratum (π_θ) and their corresponding treatment probabilities (ρ_θ). The expressions derived in Section 1 imply that any feasible parameter vector must reproduce the observed participation moments through the identities linking $(\pi_\theta, \rho_\theta)$ to $(Pr[Y=1], Pr[Y=1 \mid Z=1], Pr[Y=1 \mid Z=0], Pr[Z=1])$. For each feasible vector, I then compute the ATE and ATT directly using the closed-form expressions established earlier:

$$ATE = \pi_T - \pi_U, \qquad ATT = \frac{\pi_T \rho_T - \pi_U \rho_U}{Pr(Z=1)}.$$

To characterize the full set of parameter values consistent with the data, I explore the feasible parameter space for $\{\pi_\theta, \rho_\theta\}$ for all $\theta \in \{A, T, U, N\}$. The approach follows the logic of partial identification: instead of selecting assumptions that point-identify either estimand, I recover the entire identified set of parameter combinations implied jointly by the observed moments and the model's structure (Kline and Tamer, 2023). This proceeds as follows.

1. I construct a fine grid over all possible population shares $(\pi_A, \pi_T, \pi_U, \pi_N)$ such that $\sum_\theta \pi_\theta = 1$.

2. For each candidate vector of population shares, I use the observed moments in Table 3 to solve for the implied treatment probabilities ρ_T and ρ_U that rationalize the observed conditional participation rates.

3. I then solve for ρ_A and ρ_N using the marginal probability of victimization, $Pr(Z=1)$.

4. Any parameter vector that violates logical or probabilistic constraints (e.g., treatment probabilities outside $[0,1]$ or participation rates inconsistent with potential outcomes) is discarded.

[7] These proportions are estimated from survey data and therefore contain uncertainty. However, the large sample size yields very precise estimates, rendering the incorporation of this uncertainty into subsequent simulations negligible and immaterial to the results. For simplicity and clarity, I omit this uncertainty from the explanation.

Participation Type	N	$Pr(Z=1)$	$Pr(Y=1)$	$Pr(Y=1\|Z=1)$	$Pr(Y=1\|Z=0)$
Attended Protest	149,728	0.21 (0.001)	0.08 (0.0007)	0.14 (0.002)	0.07 (0.0007)
Community Problem-Solving	94,022	0.19 (0.001)	0.36 (0.001)	0.43 (0.004)	0.34 (0.002)
Attended Community Meeting	149,405	0.21 (0.001)	0.29 (0.001)	0.32 (0.003)	0.28 (0.001)
Attended Political Meeting	137,308	0.20 (0.001)	0.16 (0.0009)	0.18 (0.002)	0.15 (0.001)
Would vote next week	127,256	0.22 (0.001)	0.81 (0.001)	0.85 (0.002)	0.80 (0.001)

Table 3: Descriptive statistics from the LAPOP survey for each measure of political participation. N: respondents with non-missing participation and victimization responses. $Pr(Z = 1)$: victimization rate. $Pr(Y = 1)$: unconditional participation rate. $Pr(Y = 1 \mid Z = 1)$ and $Pr(Y = 1 \mid Z = 0)$: participation rates among victims and non-victims, respectively. Standard errors in parentheses.

5. For every remaining feasible parameter vector, I compute the corresponding ATE and ATT using the expressions above, as well as the implied differential selection ratio $\lambda = \rho_U/\rho_T$ and stratum-size ratio $\eta = \pi_T/\pi_U$, which jointly govern whether the ATT agrees in sign with the ATE (Section 5). Because the ATT depends on heterogeneous treatment probabilities, while the ATE depends only on population shares (Proposition 1), differences between the two estimands across feasible vectors reveal the extent to which heterogeneous exposure patterns can reverse or attenuate conclusions about victimization's political effects.

This procedure yields a distribution of ATEs and ATTs, the identified set, that fully characterizes the causal effects consistent with the LAPOP data under heterogeneous treatment probabilities. The contrast between these identified sets illustrates the empirical relevance of the theoretical results in Section 1: whenever individuals who reduce participation after victimization are more likely to be victimized (i.e., $\rho_U > \rho_T$), the ATE and ATT diverge, often in magnitude and sometimes in sign. All algebraic derivations, feasibility constraints, and computational details are reported in Appendix C.

With this set of parameters, I can compare the realized change in outcomes under actual exposure patterns to the hypothetical change under random exposure. I allow for the effect of victimization to mobilize or demobilize participation. That is, I assume $\pi_U \geq 0$.[8]

[8]However, Figure A4 in the Appendix reports the results when the effect of victimization on participation is assumed to be only increasing.

3 Results: Crime Victimization

After repeating the algorithm detailed in the previous section for each of the participatory outcomes, I obtained 12.1 million, 16 million, 20.6 million, 18.34 million, and 3.63 million feasible parameter vectors $\{\pi_\theta, \rho_\theta\}$ $\forall\theta \in \{A, T, U, N\}$, each implying a specific λ and η, for the protest attendance, political meeting attendance, community meeting attendance, community problem-solving outcomes, and willingness to vote respectively. These are an internally consistent set of values that bound the parameter's actual value. The difference in the sizes of the parameter sets depends on how informative the data is; for infrequent outcomes, fewer parameter sets are consistent with the data.

For each set, I can then calculate the corresponding ATE and ATT of victimization on each of the measures of political participation. Figure 1 plots the results. Each point in the graph corresponds to a plausible combination of π_θ, ρ_θ for a given participatory outcome. The red diagonal line that passes through the origin indicates the points where the ATE and the ATT are equal. For all the points below, the ATE is larger than the ATT. Conversely, the ATT is larger than the ATE for all points above that line. The color of the point indicates the difference between the probability of becoming a victim for individuals who become mobilized due to victimization and the probability of becoming a victim for individuals who respond by decreasing political participation in the set.

The implied ATE is almost always larger than the corresponding ATT across all measures of non-electoral political participation. When participating in a protest is the outcome, the ATE is larger than the ATT in 99% of the parameter sets. When attending political meetings is the outcome; the same holds for 96% of the parameter sets. When it is attending community meetings, the ATE is larger in 83% of the parameter sets. When participating in community problem-solving activities is the outcome, the percentage equals 64%. Further, we can see that the points above the red diagonal line correspond to the largest differences between $\rho_T - \rho_U$. For the ATT to be larger than the ATE, selection into victimization has to be strongly correlated with the treatment, making

individuals more participative.

While the majority of the ATEs compatible with the data are positive for all four measures of non-electoral participation (85% for protest attendance, 79% for political meeting attendance, 67% for community meeting attendance, and 63% for community problem solving), the converse is true for the ATTs across measures. Specifically, only 9% of the ATTs are positive for protest attendance, only 13% for political meeting attendance, 29% for community meeting attendance, and 46% for community problem-solving. Even for the most frequently reported outcome, community meeting attendance, in most cases, the ATTs compatible with the data are negative. As reported in Table 3, the larger share of positive ATEs and negative ATTs reflects the conditional probabilities of participating in these activities given victimization implied in the respondents' answers.

The results from the willingness-to-vote question, if elections were held the following week, show that electoral participation, to the degree that this is a construct-valid question, shows a markedly different pattern than non-electoral forms of engagement. While the ATE exceeds the ATT across most non-electoral participation measures, this relationship reverses for voting intention. Here, the ATE exceeds the ATT in only 1% of parameter sets. More strikingly, in 89% of parameter sets compatible with the data, the ATE of victimization on willingness to vote is *negative*, while in 95% of the sets, the ATT is positive.

Three considerations help explain this reversal. First, voting may be relatively less costly than non-electoral participation. For the marginal individual, it requires no collective organization, no sustained community presence, and no resources beyond showing up. The mechanisms that suppress costly participation after victimization may therefore not suppress voting. Second, strata may be outcome-specific. An individual can be in the *Participate if Untreated* stratum for protest but in the *Participate if Treated* stratum for voting. With roughly 8% of respondents reporting willingness to vote and 20% victimization, the population is dominated by *Always* voters for whom voting behavior is unaffected by victimization, consistent with the empirical stickiness of electoral participation (Coppock and Green, 2016). Third, the reversal is consistent with extant evidence.

Sønderskov et al. (2022), using Danish panel data, find that victimization increases turnout by 2 to 3 percentage points within individuals but that the between-individual association is large and negative. Their finding also provides direct evidence that exposure probabilities differ across response types: individuals who would increase turnout after victimization appear less likely to be victimized in the first place.

What drives the divergence between the ATE and ATT: the relative size of the mobilized and demobilized strata, or their differential exposure to victimization? Figure 2 plots the cumulative distribution of these two quantities across all parameter vectors compatible with the LAPOP data. The left panel shows $\eta = \pi_T/\pi_U$, the ratio of mobilizers to demobilizers; the right panel shows $\lambda = \rho_U/\rho_T$, the ratio of demobilizer to mobilizer victimization risk. Both panels use a log scale, with the dashed line marking the reference value of 1.

The left panel shows that across all four non-electoral outcomes, the vast majority of the identified set has $\eta > 1$, with mobilizers outnumbering demobilizers. Protest has the highest η values, with a median around 3.5; the community outcomes cluster lower, with medians between 1.7 and 2.6. The cumulative proportion at which each curve crosses the dashed line equals the share of vectors with $ATE > 0$, since $ATE = \pi_T - \pi_U$. For voting, the pattern reverses: $\eta < 1$ for most of the identified set, indicating that demobilizers outnumber mobilizers and that the ATE is predominantly negative.

The right panel tells the complementary story. For all four non-electoral outcomes, $\lambda > 1$ across the vast majority of the identified set, meaning demobilizers face higher victimization risk than mobilizers. Protest and political meetings show the strongest pattern, with λ rising above 1 early in the cumulative distribution. Voting again reverses: $\lambda < 1$ for most vectors, meaning mobilizers are more likely to be victimized than demobilizers.

The results suggest that if victimization were random, it would likely increase non-electoral participation, consistent with Bateson (2012) and the field's prevailing understanding.[9] However,

[9]Results in Bateson (2012) are based on one wave of the LAPOP data. Section A2.3 in the Appendix replicates the

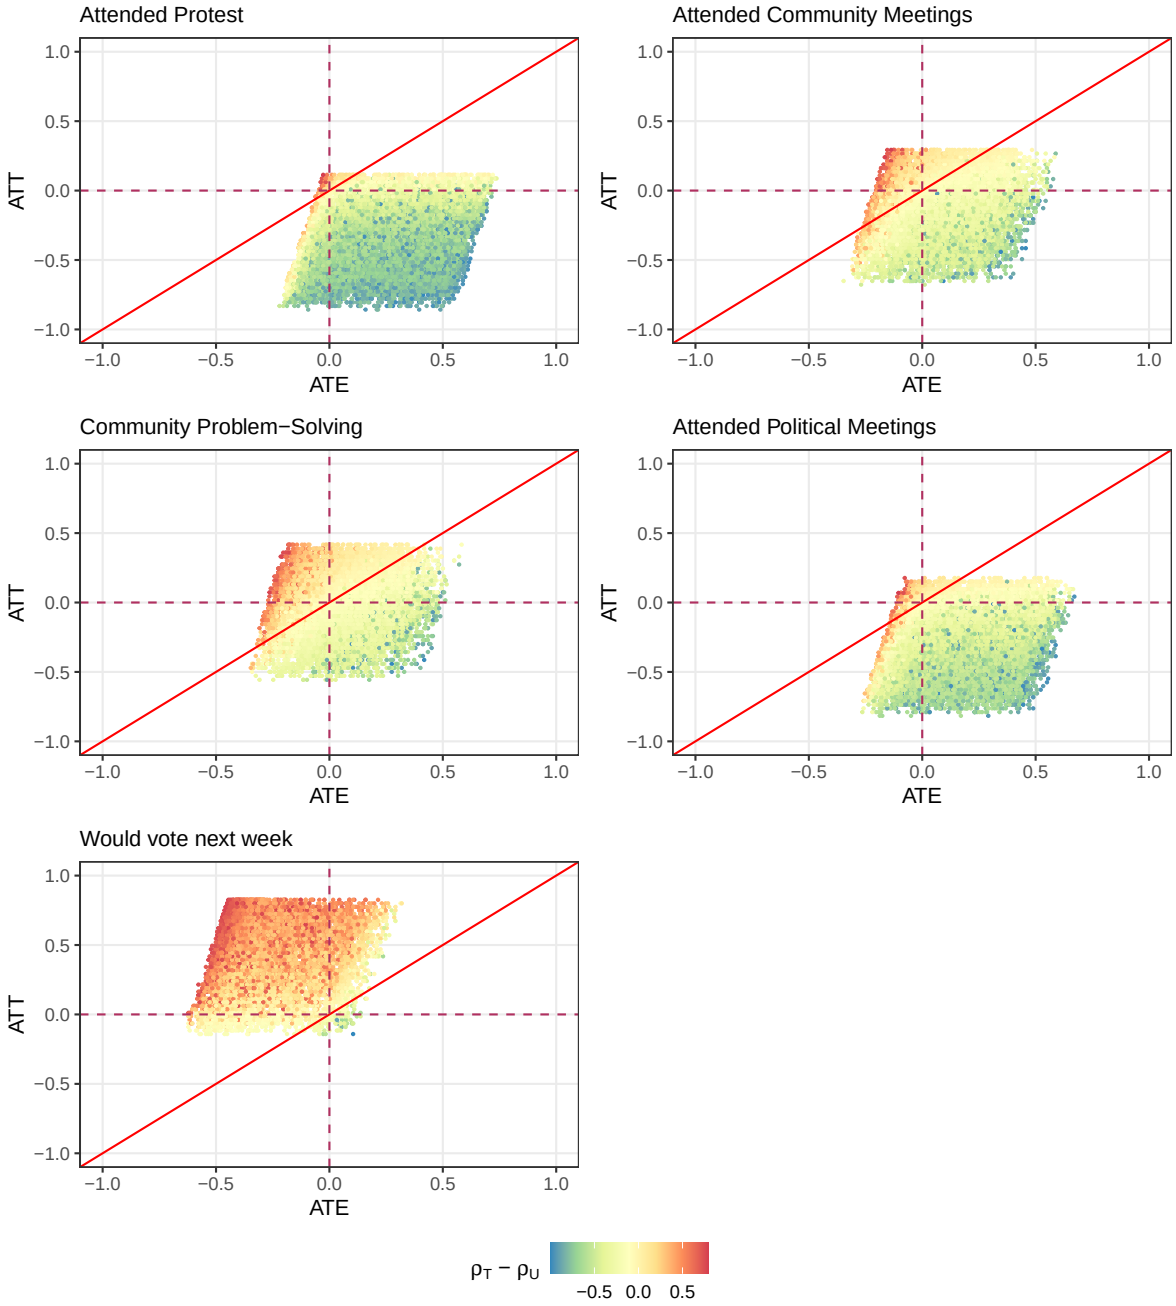

Figure 1: Figure shows the combination of potential ATTs and ATEs of victimization on four measures of political participation, conditioning on observed participation, and criminal victimization (LAPOP, 2022).

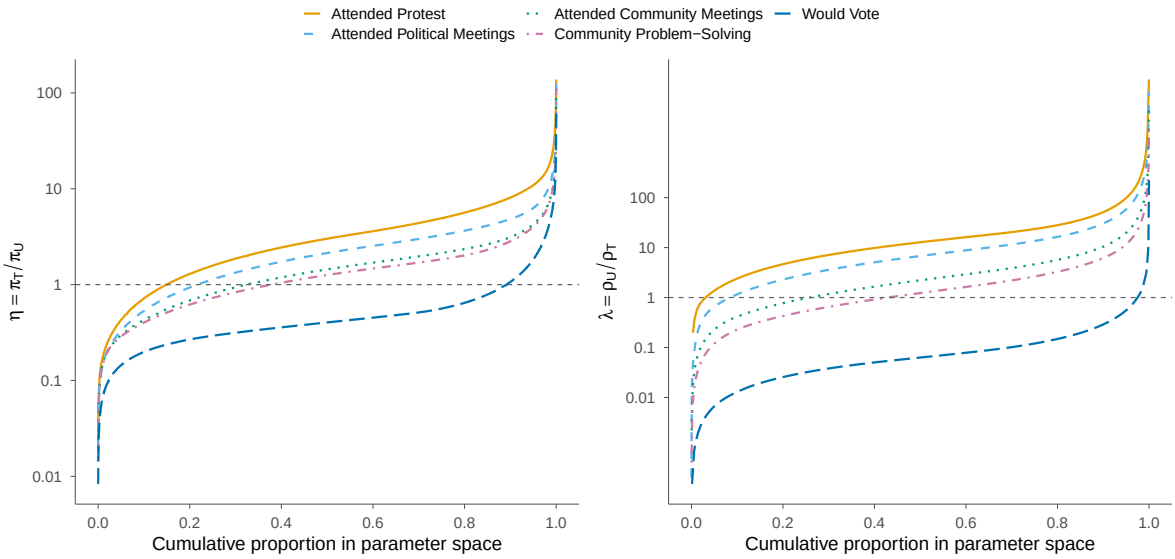

Figure 2: Cumulative distribution of the stratum-size ratio $\eta = \pi_T/\pi_U$ (left) and the differential selection ratio $\lambda = \rho_U/\rho_T$ (right) across parameter vectors compatible with the LAPOP data. Log scale. The dashed line marks 1: above it, mobilizers outnumber demobilizers (left) or demobilizers face higher victimization risk (right). The cumulative proportion at which a curve crosses 1 equals the share of the identified set with $ATE > 0$ (left) or $\rho_U > \rho_T$ (right).

because individuals who would participate less after victimization are disproportionately likely to become victims, the identified sets suggest that crime reduces all four types of non-electoral participation in the region. The pattern reverses for willingness to vote: crime, were it random, would likely reduce willingness to vote, but given actual patterns of victimization, it likely increases it.

results from Bateson (2012) with repeated cross-sections of the LAPOP data, as well as within-country analyses.

		ATE			ATT			ΔY	
Participation Type	$Pr(Z=1)$	**Min**	**Max**	$Pr(ATE<0)$	**Min**	**Max**	$Pr(ATT<0)$	**Min**	**Max**
Attended Protest	0.19	-0.23	0.76	0.15	-0.84	0.12	0.91	-0.18	0.02
Community Problem-Solving	0.19	-0.38	0.61	0.37	-0.57	0.42	0.53	-0.11	0.08
Attended Political Meetings	0.20	-0.28	0.70	0.21	-0.82	0.16	0.87	-0.17	0.03
Attended Community Meetings	0.21	-0.36	0.63	0.32	-0.68	0.30	0.71	-0.14	0.06
Would vote next week	0.22	-0.65	0.33	0.89	-0.13	0.84	0.05	-0.03	0.18

Table 4: Identified sets of ATE, ATT, and $\Delta Y = ATT \times Pr(Z=1)$ from all feasible parameter vectors compatible with the LAPOP data, without constraining ρ_U or ρ_T. Min and Max report the bounds of each identified set. $Pr(ATE<0)$ and $Pr(ATT<0)$ report the share of feasible parameter vectors for which each estimand is negative.

What are the magnitudes of the changes in participation? Recall from expression (1) that the realized change in population outcomes equals the ATT scaled by the probability of treatment. Table 4 synthesizes all the results. Participation in protests could decrease by as much as 18 pp, engagement in community problem-solving could decrease by 11 pp, political meeting attendance could decrease by 17 pp, while community meeting attendance could decrease by as much as 14 pp. Willingness to vote next week, however, at worst could decrease by 3 pp. Conversely, conditional on the data, protest attendance could *at best* increase only by 2 pp; community problem solving by 8 pp; political meeting attendance by 3 pp; and community meeting attendance by 6 pp. Hypothetical willingness to vote could increase by at most 18 pp. The *upper bound effects* within the identified set of non-electoral participation estimands are comparable in magnitude to, and often smaller than, the ATEs identified in Bateson (2012).

One concern is that these results reflect peculiarities of the LAPOP data. For example, we might be worried that LAPOP over-represents urban areas, under-represents dangerous areas, or yields survey responses that do not generalize. To address such concerns, I analyze data from Dorff (2017), who conducted a nationally representative survey of 1,000 Mexican respondents that included questions about victimization and both electoral and non-electoral participation. I apply the same partial identification procedure from the main paper to this alternative dataset. I

present the results in Table A6 in the Appendix. I find strikingly similar results that point in the same direction: the identified sets suggest that while the ATE is likely positive, the ATT is almost certainly smaller and likely negative.

4 Panel Evidence from Mexico

The previous section established that demobilization among the victimized is consistent with the majority of the parameter space arising from pooled data from 20 Latin American countries. The natural next question is where in that space we actually are. Answering it requires evidence on who is selected into victimization and how they respond, which cross-sectional data cannot provide. I turn to a panel survey from Mexico, one of the countries in the LAPOP sample, that follows the same individuals over time.

I use a quarterly rotating panel survey of Mexican city residents that asks questions related to crime and insecurity, the National Survey of Urban Public Safety, ENSU, for its name in Spanish (INEGI, 2024). This survey has been conducted four times per year since 2013[10] and, starting in 2017, has included direct questions regarding personal victimization. This rolling survey panel allows me to observe individuals at five different points over 15 months. In each quarter, respondents are asked about experiences with the state, crime, and conflict, and in at least two waves per year, they are asked about direct victimization. I keep all survey waves from 2017, when crime victimization was first incorporated, through the last quarter of 2023, resulting in 560k responses from 174K unique respondents collected over 27 quarters.

Unlike the LAPOP data, the panel structure of ENSU allows me to provide direct evidence on both conditions that the framework requires. First, by comparing future victims with non-victims before any victimization occurs, I assess whether exposure is selective: if it is, future victims should differ systematically from their neighbors even before becoming victims. Second, by exploiting within-individual variation, I estimate how victimization changes attitudes and be-

[10]Except the second quarter of 2020 due to the COVID emergency.

havior among those who experience it, providing evidence on the direction of response among the treated.[11]

To assess whether exposure to victimization is selective, I identify ENSU questions that map onto the characteristics the framework calls for to vary across strata: mobility patterns, exposure to criminal and violent behavior, feelings of unsafety, and interpersonal conflicts. Higher mobility increases exposure to crime (Rader, 2004; Browning, Pinchak and Calder, 2021); witnessing criminal behavior directly indexes victimization risk; and fear of crime correlates with perceived vulnerability (Guedes, Domingos and Cardoso, 2018). Critically, these same characteristics are also correlates of political participation: extroverted, mobile, community-oriented, and risk-prone individuals participate more (Gerber et al., 2011; Kam, 2012; Ley, 2022). While ENSU does not directly measure political participation, these behavioral and attitudinal measures serve as proxies for the participatory profiles that distinguish strata in the framework.

To assess reactions to victimization, I estimate how victimization changes these same attitudes and behaviors among those who experience it. I exploit within-individual variation using two estimators: a two-way fixed effects estimator and the fixed-effect counterfactual estimator proposed by Liu, Wang and Xu (2022). I focus on measures of crime prevention behavior, fear of crime, trust in law enforcement, and mobility. If victimization demobilizes, victims should become less mobile, more fearful, and less trusting; the behavioral profile of the *Participate if Untreated* stratum.

4.1 Exposure to Crime

I subset the data and retain only respondents who never reported criminal victimization and respondents who reported criminal victimization in some, but not all, survey waves. For the latter, I retain responses only before their first reported instance of victimization. Consequently, I compare the responses of this group of future victims *before they reported victimization* with those of their

[11] Like any survey measure, self-reported victimization may be subject to heterogeneous reporting: individuals may interpret events as criminal or non-criminal differently, and such sensitivity could correlate with political attitudes and behavior (Skogan, 1982; Boulding, Mullenax and Schauer, 2022). Two descriptive analyses in Appendix Appendix E show that self-reported victimization tracks municipal-level criminal victimization rates as expected.

never-victimized neighbors—either from the same sampling unit or the same locality—in the same quarter.

I estimate

$$PM_{i[q,g]} = \beta_1 * FutureVictim_{i[q,g]} + \theta_{q \times g} + \epsilon_{i[q,g]} \tag{3}$$

where $PM_{i[q,g]}$ is respondent $i's$ self-reported behavior, experiences, or attitude in year-quarter q, and neighborhood/census tract or locality g, or a binary measure that takes the value of 1 if respondent i reports that type of behavior or experience and 0 otherwise. $FutureVictim_{i[q,g]}$ is a binary variable that takes the value of 1 when respondent i will report having been a victim of crime in a future survey and 0 otherwise. $\theta_{q \times g}$ represent either $Neighborhood \times Quarter$ fixed effects, or $Locality \times Quarter$ fixed effects. $\epsilon_{i[q,g]}$ are robust errors clustered at the primary sampling unit.

Future victims differ systematically from non-victims before victimization occurs (Figure 3). When we compare future victims with never-victims living in the same locality in that same quarter (dark blue) or the same small neighborhood (light blue), future victims report having more conflicts, including conflicts that resulted in physical violence, conflicts with authorities, conflicts with the police, and conflicts with strangers on the street, *before* being victimized. Future victims also witnessed gangs, robberies, and vandalism, and heard gunshots more frequently. They also report feeling unsafe in the city, on the streets, on public transport, and even in their house. Last, they report leaving their home more frequently.

Overall, the results link future victimization with increased exposure to crime and show that victims and non-victims were differentially exposed to crime even before becoming victims. These findings are consistent with future victims disproportionately resembling the *Participate if Untreated* stratum: before victimization, they are more mobile, more exposed to criminal activity, and more embedded in conflictual social interactions, which matches the behavioral profile associated with higher political participation in extant research (Gerber et al., 2011; Kam, 2012).

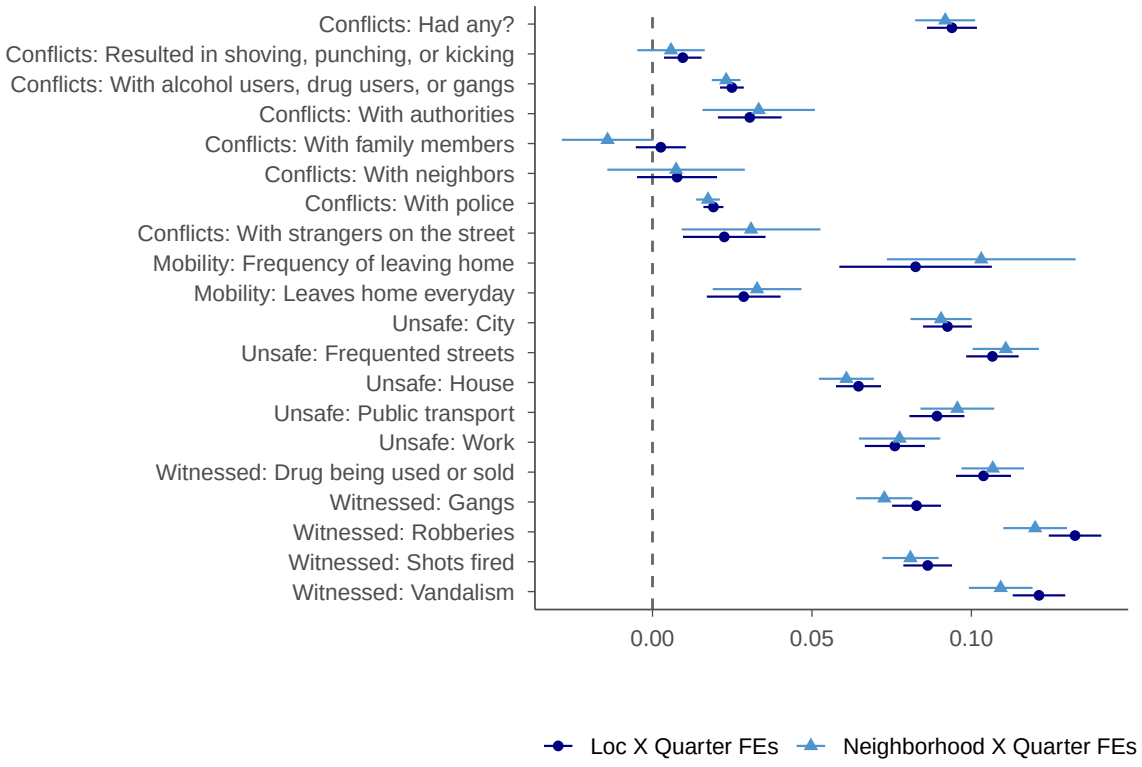

Figure 3: Cross-sectional differences between future victims and non-victims living in the same locality (dark blue) and the same neighborhood/census tract (light blue) on self-reported outcomes relating to exposure to crime. Robust standard errors clustered at the primary sampling unit. Full regression results in Table A12.

4.2 Reactions to Crime

I next estimate how victimization changes attitudes and behavior among those who experience it.

$$PM_{i[q]} = \beta_1 * Victim_{i[q]} + \theta_i + \gamma_q + \epsilon_{i[q]} \tag{4}$$

where $PM_{i[q]}$ is either a continuous measure of respondent $i's$ self-reported behavior, experiences, or attitude in year-quarter q, and neighborhood or locality g, or a binary measure that takes the value of 1 if respondent i reports that type of behavior, experience or attitude and 0 otherwise. $Victim_{i[q]}$ is an indicator variable that takes the value of 1 when respondent i reports having been victimized in the past six months in quarter q and 0 otherwise. Because victimization is asked every six months, while response variables are asked every three months, each respondent's victimization report is extended to cover the prior period. θ_i are individual fixed effects while γ_q are period fixed effects. $\epsilon_{i[q]}$ are robust errors clustered at the primary sampling unit. I use a TWFE estimator of the ATT as well as the counterfactual estimator with equal weights proposed by Liu, Wang and Xu (2022).

Victimization changes attitudes and behaviors in a way consistent with demobilization, regardless of the preferred estimator (Figure 4). Compared with their behavior before becoming crime victims, respondents report leaving home less frequently, visiting friends and family less often, and taking more steps to prevent crime. They become less mobile and less likely to say they leave home every day. They trust law enforcement less, fear crime more, and adopt more constrained behaviors.

These behavioral and attitudinal changes are mechanisms through which demobilization is theorized to operate. Reduced mobility and social withdrawal limit the interpersonal interactions that facilitate political engagement (Gerber et al., 2011; Kam, 2012). Increased fear raises the perceived costs of public engagement (Ley, 2018; Skogan, 1986). ENSU does not measure political participation directly, but the within-individual changes it documents consistently move in the

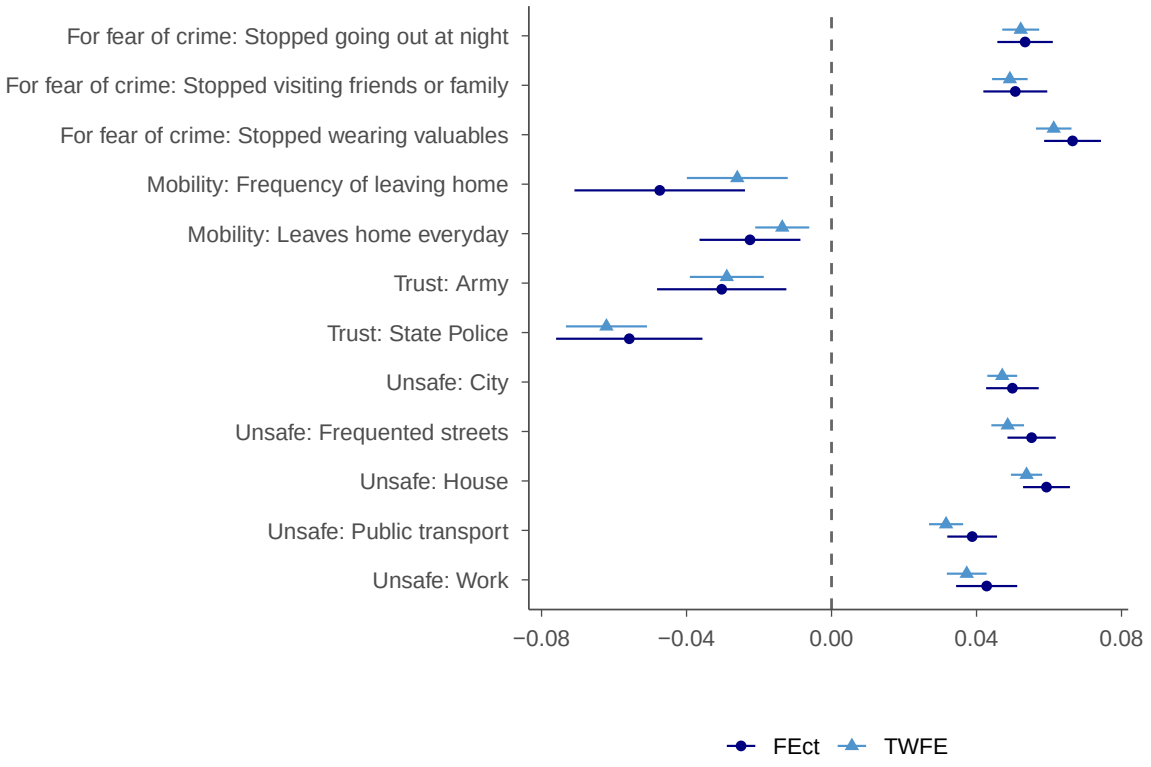

Figure 4: Within-individual effects of victimization on self-reported behavioral and attitudinal measures, estimated with the fixed-effect counterfactual estimator from Liu, Wang and Xu (2022) (dark blue) and a two-way fixed effects estimator (light blue). Robust standard errors clustered at the primary sampling unit. Full regression results in Table A13.

direction that the participation literature predicts should reduce engagement.

5 Disciplining the Identified Set with Panel Evidence

The LAPOP cross-section mapped the range of plausible effects in a pooled sample from Latin American countries. The ENSU panel showed that, in Mexico, individuals with the behavioral profile of demobilizers are more exposed to crime, and that victimization moves behavior in the demobilizing direction. If this pattern generalizes across the region, the remaining question is whether mobilizers are numerous enough relative to demobilizers to keep the ATT positive despite their lower exposure.

To address this question, recall from Section 1 that whether the ATT agrees in sign with the ATE depends on two quantities:

1. The *differential selection ratio*, $\lambda = \rho_U/\rho_T$: how much more likely demobilizers are to be victimized compared to mobilizers. When $\lambda = 1$, both strata face identical victimization probabilities and sign agreement is guaranteed (Corollary 2). When $\lambda > 1$, demobilizers are disproportionately exposed to treatment.

2. The *stratum-size ratio*, $\eta = \pi_T/\pi_U$: how numerous mobilizers are relative to demobilizers. When $\eta > 1$, mobilizers outnumber demobilizers in the population.

When $ATE > 0$, Proposition 3 establishes that the ATT is negative if and only if $\lambda > \eta$. That is, the ATT reverses sign relative to the ATE when the overexposure of demobilizers exceeds their relative scarcity. The ENSU panel showed that in Mexico, future victims are systematically more exposed than non-victims before victimization, and that victimization moves behavior in the demobilizing direction within individuals. Together, these findings support the conclusion that $\lambda > 1$ for Mexico: demobilizers likely face a higher risk of victimization than mobilizers.

I assume this directional condition holds across Latin America, without imposing any additional assumptions on the magnitude of λ. Given $\lambda > 1$, the sign of the ATT depends entirely on

η, the ratio of mobilizers to demobilizers in the population. Since $ATE = \pi_T - \pi_U$, when the ATE is small, the two groups are of comparable size and η is close to 1. Published ATE estimates for non-electoral participation are on the order of 1 to 4 percentage points (Bateson, 2012). Small ATEs imply that mobilizers do not vastly outnumber demobilizers. These published estimates are point-identified under conditional ignorability and serve here as an informative prior over the ATE's identified set, not as a substitute for it. The question is whether this modest imbalance is enough to sustain a positive ATT when demobilizers are overexposed.

Table 5 shows how the ATT sign depends on restrictions to η. I fix $\lambda > 1$ and gradually decrease the maximum number of mobilizers per demobilizer. For each restriction, among LAPOP parameter vectors with $ATE > 0$ and within the bound, I compute the proportion with $ATT < 0$. When this proportion hits 100%, the ATT is *signed*: all feasible vectors agree on direction.

	Bound on $\eta = \pi_T/\pi_U$ (mobilizers per demobilizer)				
	Unrestricted	≤ 2	≤ 1.5	≤ 1.1	≤ 1.04
Share of feasible vectors with $ATT < 0$, given $ATE > 0$ and $\lambda > 1$.					
Protest	0.90	0.98	0.99	>0.99	**Signed** $-$
Community Problem-Solving	0.71	0.80	0.87	0.97	>0.99
Community Meetings	0.79	0.89	0.93	0.97	0.98
Political Meetings	0.89	0.97	0.98	>0.99	**Signed** $-$

Table 5: Sensitivity of ATT sign identification to restrictions on the stratum-size ratio. Cells report the share of LAPOP-compatible parameter vectors with $ATT < 0$, conditional on $ATE > 0$ and $\lambda > 1$ (demobilizers more exposed). Columns impose progressively tighter upper bounds on $\eta = \pi_T/\pi_U$. **Signed** $-$: 100% of feasible vectors yield $ATT < 0$.

Even without restricting η, imposing only $\lambda > 1$, between 71% and 90% of parameter vectors yield a negative ATT across the four non-electoral outcomes. At $\eta \leq 1.5$, all four outcomes exceed 87%. At $\eta \leq 1.1$, protest and political meetings exceed 99%. At $\eta \leq 1.04$, protest and political

meetings are signed negatively: every feasible parameter vector agrees on the direction. Since published ATE estimates are small, on the order of 1 to 4 percentage points, mobilizers cannot vastly outnumber demobilizers, placing η well within this range. Given the magnitude of published effects and the directional evidence from the ENSU panel, imposing $\lambda > 1$ signs the effect among the victimized negative for protest and political meetings, and negative for over 97% of the feasible parameter space for community meetings and community problem-solving. Overall, results show that the mobilization finding for non-electoral participation does not hold up when shifting from the population average to those actually victimized. A corresponding analysis for voting, where the sign-flip runs in the opposite direction, is reported in Appendix A4.2.

6 Applications to other Forms of Violence

Criminal victimization is not the only form of violence where selective exposure and heterogeneous responses plausibly coincide. I apply the partial identification procedure to civil war abduction in Uganda and police contact in the United States. Unlike the crime application, I cannot constrain the identified sets with panel evidence in these settings. The extensions therefore cannot establish sign reversal with the same force. What they can establish is that the identified sets are consistent with the framework's predictions: in both settings, the region of the parameter space where the ATE and ATT disagree in sign is large, and the direction of disagreement mirrors the crime application.

6.1 Civil War Violence

Studies of civil war violence have generally found that exposure makes victims more likely to participate in politics, attributing this to post-traumatic growth and increased prosociality (Bellows and Miguel, 2009; Blattman, 2009; Voors et al., 2012). One influential paper is Blattman (2009), which examines how forced recruitment into the LRA insurgency in Northern Uganda shaped downstream political participation. Leveraging theorized quasi-random variation in abduction, he estimates the ATE and finds that abduction significantly increased voting and doubled the likelihood of becoming a community leader.

I apply the partial identification procedure to data from Blattman (2009) to derive the identified sets of ATEs, ATTs, and ρs. Detailed results are reported in Section A6.1. For electoral participation, 77% of the ATEs in the identified set are positive, and 75% of the ATTs are as well — unlike the LAPOP results, the two estimands largely agree. For non-electoral participation, however, the pattern mirrors the crime application. Only 35% of the identified ATEs imply that abduction increased community leadership, despite the original point estimate being positive and significant. And between 93%–96% of the identified ATTs are negative. If abduction were random, it would likely increase voting and some forms of civic engagement. If it were selectively experienced, the identified set suggests it may have depressed non-electoral participation among the abducted.

If abduction was random, as theorized by the author, then by Corollary 1, the ATE equals the ATT, and research designs targeting either estimand are equally suitable for measuring actual changes in participation. The partial identification exercise shows what is at stake when that assumption fails: if exposure to abduction was even moderately selective, the positive effects on non-electoral participation that the ATE implies may not describe what happened to those who were actually abducted. The partial identification exercise does not adjudicate whether abduction was random, but it clarifies the stakes of that assumption for the substantive conclusion.

6.2 Contact with the Police

Research on the political consequences of police contact presents competing findings. Some work documents that contact teaches individuals they are politically marginal, suppressing engagement (Weaver and Lerman, 2010; Soss and Weaver, 2017). Others find that police contact, whether direct or indirect, stimulates participation (Walker, 2014, 2020). That police engage in statistical discrimination is well established (Ba et al., 2021; Mummolo, 2017), making this a setting where exposure is plausibly selective.

Walker (2020) finds that direct and indirect police contact is associated with higher rates of both electoral and non-electoral participation, targeting the ATE as the primary estimand. I apply

the partial identification procedure to one of the data sources used, the National Crime and Politics Survey, a nationally representative U.S. survey conducted in 2013. I focus on direct police contact and use all available measures of electoral and non-electoral participation. Results are reported in Table A11.

The pattern mirrors the crime application. While most parameter sets show positive ATEs for non-electoral participation, the ATTs are predominantly negative. For voting, the pattern again reverses: 89% of parameter sets yield negative ATEs, while only 7% show negative ATTs. If police contact were random, it would increase non-electoral participation. If it is selectively experienced in ways that coincide with heterogeneity in reactions, as the abundant literature suggests it is, the identified set suggests it depresses non-electoral engagement among those most frequently stopped while increasing their electoral participation.

7 Guidance for Researchers

The results in this paper call into question substantive conclusions about the participatory consequences of violence drawn from population-average effects. In this section, I distill the insights from the framework and empirical findings into a checklist for empirical researchers, applicable beyond violence.

1. What estimand does the analysis target? The estimand links theory to empirics by specifying the population over which a causal effect is defined and, therefore, the counterfactual comparison it supports. Researchers should understand which population their estimate is defined over and why that population is the relevant one for the claim.

2. Does the estimand incorporate selective exposure in a manner consistent with the theory? In specifying the population, the estimand also determines how heterogeneous treatment probabilities are incorporated into the empirical analysis. The ATE weights response types by their population shares; the ATT weights them by their exposure probabilities. Both are useful quantities. The ATE answers how treatment would change outcomes if exposure were equal across

subpopulations. The ATT answers how treatment changes outcomes given who is actually exposed. When these diverge, each supports a different substantive claim. Researchers should be explicit about which question they are asking and whether their estimand answers it.

3. Are the findings you compare defined for the same population? When engaging with the empirical literature, researchers should consider whether the studies they compare define effects for the same population. In the presence of selective exposure, apparent disagreement across studies may reflect differences in the estimands being targeted rather than in mechanisms or contexts. Making the estimand explicit clarifies whether findings are commensurable, which literature the study contributes to, and what the prior expectation for effects should be.

4. If treatment and outcome are binary, characterize what the data allow. The partial identification procedure requires four quantities: $Pr(Z = 1)$, $Pr(Y = 1)$, $Pr(Y = 1|Z = 1)$, and $Pr(Y = 1|Z = 0)$. From these, the researcher can characterize the identified set of ATEs and ATTs consistent with the observed moments. Any treatment and outcome can be made binary — continuous outcomes can be dichotomized at a meaningful threshold, and ordinal scales can be collapsed — though this comes at some loss of information. The binary-binary setting is thus widely applicable. Computational details and replication code are provided in Appendix C.

5. Interpret the identified set. For each estimand, check whether all feasible parameter vectors yield the same sign. Table 2 provides a quick diagnostic: given an estimated ATE and a belief about how selective exposure is, researchers can assess whether sign reversal is possible. If yes, the sign is identified. If both the ATE and ATT are sign-identified and agree, the conclusion is robust to selective exposure. If they disagree, the substantive claim supported by one estimand may not hold for the other. When the identified set for a given estimand spans both signs, the distribution across quadrants and the implied selection parameters ($\rho_T - \rho_U$) indicate what patterns of selective exposure would be required to sustain each conclusion. Theoretical knowledge of the setting helps assess which regions of the parameter space are plausible and which require difficult-to-justify

selections. When panel or auxiliary data are available, point estimates can be located within the identified set to narrow the plausible range.

6. Report what the exercise reveals about the estimand you targeted. If the identified set confirms that the ATE and ATT share the same sign, report this as evidence that the conclusion is robust to selective exposure. If they diverge, report the identified set and discuss what it implies for the subpopulation the estimate is defined over. When there are strong theoretical or empirical reasons to believe exposure is random, the exercise serves a different purpose: it clarifies the epistemological stakes of that assumption by showing what conclusions would follow if it fails. In either case, the exercise makes explicit what the data are and are not consistent with, and allows the researcher and the reader to assess how much the substantive conclusion depends on the choice of estimand.

8 Conclusion

Most of what we know about the political consequences of violence from empirical causal research comes from counterfactuals in which everyone is exposed to a treatment that falls on a select few. In this paper, I showed that the population-average effect of violence on political participation and the effect among the exposed can yield opposite answers. The literature on criminal victimization has read positive population-average effects on non-electoral participation as evidence that violence produces a more engaged citizenry. Yet I find that those who become victims of crime were already more exposed than their neighbors before victimization, and that they react to it by becoming less mobile, more fearful, and less trusting of law enforcement. These are the behavioral changes that the participation literature associates with withdrawal from political life. Violence could well mobilize on average if everyone were equally exposed to it. But the victimized are often the vulnerable, and among them, violence is more likely than not to demobilize.

That targeting is selective and that the vulnerable are disproportionately exposed is well established in the theoretical literature on violence. What has not been established is the consequence

for empirical causal research: when these features coincide, the population over which an effect is defined determines its sign. For criminal victimization in Latin America, where panel evidence supports this conclusion, and plausibly for civil war abduction in Uganda and police contact in the United States, heterogeneous responses to violence coincide with selective exposure. A better understanding of who is exposed to violence and how they differ from those who are not should inform not only the design of empirical studies but also the interpretation of the estimates they produce.

Whether violence erodes or strengthens civic engagement among those who experience it determines whether affected communities can hold their governments accountable. If the victimized withdraw from political life while the population average suggests mobilization, the political system registers the preferences of those least affected by violence, not those most in need of a governmental response. This is true not only of violence but of any treatment that falls selectively and that we would not seek to universalize. For this class of treatments, the effect among the exposed is the quantity of first-order interest.

References

Ba, Bocar A., Dean Knox, Jonathan Mummolo and Roman Rivera. 2021. "The Role of Officer Race and Gender in Police-Civilian Interactions in Chicago." *Science* 371(6530):696–702.

Bateson, Regina. 2012. "Crime Victimization and Political Participation." *American Political Science Review* 106(3):570–587.

Bellows, John and Edward Miguel. 2009. "War and Local Collective Action in Sierra Leone." *Journal of Public Economics* 93(11):1144–1157.

Blanco, Luisa R. 2013. "The Impact of Crime on Trust in Institutions in Mexico." *European Journal of Political Economy* 32:38–55.

Blattman, Christopher. 2009. "From Violence to Voting: War and Political Participation in Uganda." *American Political Science Review* 103(2):231–247.

Boggs, Sarah L. 1965. "Urban Crime Patterns." *American Sociological Review* 30(6):899–908.

Boulding, Carew, Shawnna Mullenax and Kathryn Schauer. 2022. "Crime, Violence, and Political Participation." *International Journal of Public Opinion Research* 34(1):edab032.

Browning, Christopher R., Nicolo P. Pinchak and Catherine A. Calder. 2021. "Human Mobility and Crime: Theoretical Approaches and Novel Data Collection Strategies." *Annual Review of Criminology* 4(Volume 4, 2021):99–123.

Coppock, Alexander and Donald P. Green. 2016. "Is Voting Habit Forming? New Evidence from Experiments and Regression Discontinuities." *American Journal of Political Science* 60(4):1044–1062.
URL: *https://onlinelibrary.wiley.com/doi/abs/10.1111/ajps.12210*

Dorff, Cassy. 2017. "Violence, Kinship Networks, and Political Resilience: Evidence from Mexico." *Journal of Peace Research* 54(4):558–573.

Frangakis, Constantine E. and Donald B. Rubin. 2002. "Principal Stratification in Causal Inference." *Biometrics* 58(1):21–29.

Gerber, Alan S., Gregory A. Huber, David Doherty, Conor M. Dowling, Connor Raso and Shang E. Ha. 2011. "Personality Traits and Participation in Political Processes." *The Journal of Politics* 73(3):692–706.

Gilligan, Michael J., Benjamin J. Pasquale and Cyrus Samii. 2014. "Civil War and Social Cohesion: Lab-in-the-Field Evidence from Nepal." *American Journal of Political Science* 58(3):604–619.

Guedes, Inês Maria Ermida Sousa, Sofia Patrícia Almeida Domingos and Carla Sofia Cardoso. 2018. "Fear of Crime, Personality and Trait Emotions: An Empirical Study." *European Journal of Criminology* 15(6):658–679.

Hale, C. 1996. "Fear of Crime: A Review of the Literature." *International Review of Victimology* 4(2):79–150.

Heckman, James J. and Edward Vytlacil. 2005. "Structural Equations, Treatment Effects, and Econometric Policy Evaluation." *Econometrica* 73(3):669–738.

INEGI, Instituto Nacional de Estadística y Geografía. 2024. "Encuesta Nacional de Seguridad Pública Urbana (ENSU).".

Kalyvas, Stathis N., Peter Lange, Robert H. Bates, Ellen Comisso, Peter Hall, Joel S. Migdal and Helen V. Milner. 2006. *Logic of Violence in Civil War*. Cambridge University Press.

Kam, Cindy D. 2012. "Risk Attitudes and Political Participation." *American Journal of Political Science* 56(4):817–836.

Kline, Brendan and Elie Tamer. 2023. "Recent Developments in Partial Identification." *Annual Review of Economics* 15:125–150.

URL: *https://www.annualreviews.org/content/journals/10.1146/annurev-economics-051520-021124*

Knox, Dean, Will Lowe and Jonathan Mummolo. 2020. "Administrative Records Mask Racially Biased Policing." *American Political Science Review* 114(3):619–637.

Koos, Carlo and Richard Traunmüller. 2024. "The Gendered Costs of Stigma: How Experiences of Conflict-Related Sexual Violence Affect Civic Engagement for Women and Men." *American Journal of Political Science* n/a(n/a).

LAPOP, Lab. 2022. "The AmericasBarometer." www.vanderbilt.edu/lapop.

Ley, Sandra. 2018. "To Vote or Not to Vote: How Criminal Violence Shapes Electoral Participation." *Journal of Conflict Resolution* 62(9):1963–1990.

Ley, Sandra. 2022. "High-Risk Participation: Demanding Peace and Justice amid Criminal Violence." *Journal of Peace Research* 59(6):794–809.

Liu, Licheng, Ye Wang and Yiqing Xu. 2022. "A Practical Guide to Counterfactual Estimators for Causal Inference with Time-Series Cross-Sectional Data." *American Journal of Political Science* 68(1):160–176.

Lundberg, Ian, Rebecca Johnson and Brandon M. Stewart. 2021. "What Is Your Estimand? Defining the Target Quantity Connects Statistical Evidence to Theory." *American Sociological Review* 86(3):532–565. Publisher Copyright: © American Sociological Association 2021.

Marshall, John. 2022. "Tuning In, Voting Out: News Consumption Cycles, Homicides, and Electoral Accountability in Mexico.".

Mesquita, Ethan Bueno De and Scott A. Tyson. 2020. "The Commensurability Problem: Conceptual Difficulties in Estimating the Effect of Behavior on Behavior." *American Political Science Review* 114(2):375–391.

Mummolo, Jonathan. 2017. "Modern Police Tactics, Police-Citizen Interactions, and the Prospect for Reform." *Journal of Politics* 80(1):1–15.

Rader, Nicole. 2004. "The Threat of Victimization: A Theoretical Reconceptualization of Fear of Crime." *Sociological Spectrum* 24(6):689–704.

Samii, Cyrus, Laura Paler and Sarah Zukerman Daly. 2017. "Retrospective Causal Inference with Machine Learning Ensembles: An Application to Anti-recidivism Policies in Colombia." *Political Analysis* 24(4):434–456.

Skogan, Wesley. 1986. "Fear of Crime and Neighborhood Change." *Crime and Justice* 8:203–229.

Skogan, Wesley G. 1982. Methodological Issues in the Measurement of Crime. In *The Victim in International Perspective: Papers and Essays Given at the "Third International Symposium on Victimology" 1979 in Münster/Westfalia*, ed. Hans Joachim Schneider. de Gruyter.

Skogan, Wesley G. 1995. "Crime and the Racial Fears of White Americans." *The Annals of the American Academy of Political and Social Science* 539:59–71.

Slough, Tara and Scott A. Tyson. 2023. "External Validity and Meta-Analysis." *American Journal of Political Science* 67(2):440–455.

Sønderskov, Kim Mannemar, Peter Thisted Dinesen, Steven E. Finkel and Kasper M. Hansen. 2022. "Crime Victimization Increases Turnout: Evidence from Individual-Level Administrative Panel Data." *British Journal of Political Science* 52(1):399–407.

Soss, Joe and Vesla Weaver. 2017. "Police Are Our Government: Politics, Political Science, and the Policing of Race–Class Subjugated Communities." *Annual Review of Political Science* 20(1):565–591.

Vilalta, Carlos. 2020. "Violence in Latin America: An Overview of Research and Issues." *Annual Review of Sociology* 46(1):693–706.

Visconti, Giancarlo. 2020. "Policy Preferences after Crime Victimization: Panel and Survey Evidence from Latin America." *British Journal of Political Science* 50(4):1481–1495.

Voors, Maarten J., Eleonora E. M. Nillesen, Philip Verwimp, Erwin H. Bulte, Robert Lensink and Daan P. Van Soest. 2012. "Violent Conflict and Behavior: A Field Experiment in Burundi." *American Economic Review* 102(2):941–964.

Walker, Hannah L. 2014. "Extending the Effects of the Carceral State: Proximal Contact, Political Participation, and Race." *Political Research Quarterly* 67(4):809–822.

Walker, Hannah L. 2020. "Targeted: The Mobilizing Effect of Perceptions of Unfair Policing Practices." *The Journal of Politics* 82(1):119–134.

Weaver, Vesla M. and Amy E. Lerman. 2010. "Political Consequences of the Carceral State." *American Political Science Review* 104(4):817–833.